

COMP 429/529- Parallel Programming: Assignment 3

Due: 11.59 pm June 5, 2020

Notes: You may discuss the problems with your peers but the submitted work must be your own work. Please submit your report and source code through blackboard. This assignment is worth 20% of your total grade.

Iterative Sparse Matrix-Vector Multiplication (iSpMV)

In this assignment you will implement the sparse matrix-vector multiplication discussed in class (Lecture 23) using MPI. Along with this project description, studying Lectures 22 and 23 might be very helpful for the assignment. Since implementing/debugging parallel codes under message passing can be more time consuming than under threads, give yourself sufficient time to complete the assignment. In addition, this assignment requires you to conduct various performance studies. Please also allow yourself enough time for performance measurements. The performance results will be collected on the KUACC cluster. For instructions about how to compile and run MPI jobs on KUACC, please read the Environment section carefully.

Background

SpMV plays an important role in a variety of scientific and engineering applications. Iterative version of SpMV (iSpMV) is a key operation in many graph-based data mining and machine learning algorithms. With the emerge of big data, the matrices can be so large that the SpMV has to be executed in parallel. In fact, iSpMV is the backbone of Page Ranking algorithm used by Google to rank pages for the search engine. This type of kernels involves multiplication of a matrix having sparsely-located non-zero elements with a vector. An entry in the matrix indicates a link from a page to another. For example, if the matrix entry (i, j) is non-zero, that means there is a link from webpage i to webpage j . Since there is no link from every page to every other page, matrix contains a lot of zero entries.

An SpMV is typically performed in the following form, $Ax = y$, such that A is a sparse matrix of size $M \times N$, x is a vector of size N , and y is a vector of size M . In the iterative version of SpMV, input vector is also the output vector, so we are solving $Ax = x$ in a time loop.

There are different types of storage formats available for sparse matrices. Since storing zeros wastes a lot of memory space, we only store the non-zero elements and use metadata to identify the location of non-zero entries in the matrix. The most common format to store sparse matrices is compressed storage row (CSR), which is shown in below figure. The *val* array contains matrix entries, the *ind* array refers to the column indices, and the *ptr* array refers to the rows.




Figure 1: CSR format - an example sparse matrix

Serial Code

We are providing you with a working serial program that implements iSpMV. The provided code reads from a file containing sparse matrix from the Florida sparse matrix collection (source: <https://sparse.tamu.edu/>). It then multiplies the read matrix with a vector iteratively in a time steps loop.

Running The Code

- The serial code that we have provided includes a Makefile that is needed to compile the serial code. To compile and run the code, execute the following commands.

```

1 To compile, type
2     make
3 To clean the objects and executables
4     make clean
5 Example run:
6     ./build/spmv <input-file>.mtx <number-of-time-steps>

```

- You could use the campus clusters or your local machines to develop and test your implementations. If you want to work locally, you need to install an MPI library.

Parallel Code

In this assignment, you are expected to parallelize the given serial SpMV code in three different tasks, one of them being optional. In the first task, you have to implement row-wise parallelization for the matrix-vector multiplication (Lecture 23 page 33). The second task, you are asked to introduce multithreaded parallelism with OpenMP on top of MPI for the row-wise partitioning. The third task is optional and you are expected to implement load-balanced partitioning to the same problem (Lecture 23 page 31).

Part I

Row-wise Partitioning and Parallelization with MPI: In this part of the assignment, you have to parallelize the given SpMV serial code by using only MPI. In parallelizing the code, the sparse matrix is to be partitioned by dividing it by rows. This parallelization

strategy is illustrated in Figure 2.




Figure 2: Row-wise parallelization of SpMV with MPI

As shown in the figure, a sparse matrix of size $M \times N$ is partitioned across P processes, so that each process has a matrix chunk with size $(M/P) \times N$, while the multiplier vector of size N is the same for all processes. Each of these processes performs matrix-vector multiplication on its matrix chunk. After multiplying all rows in the matrix chunk with the vector, each process will have its own vector of results with size (M/P) . To generate the final output, these vectors of results are collected from all processes and appended to produce the final vector of size M in process 0. In the iterative version of iSpMV, the output vector should be distributed to all processors to update their vector x for the next iteration.

Part II

Part I with MPI+OpenMP: In this part of the assignment, you have to introduce multi-threaded parallelism to each process in the row-wise parallelization code from the first task. You are expected to implement the multithreaded parallelism using OpenMP. Parallelize across the rows assigned to a single process.

Part III - Optional

Load Balancing - Bonus: The row-wise partitioning partitions the rows equally among processes. However, depending on the sparsity pattern of the matrix, the number of non-zeros that each process will compute may not be balanced. In other words, equal number of rows per process does not guarantee that each process will be responsible for the same amount of workload. Instead, in this part, you will still partition row-wise but considering the number of non-zeros per process. That means each process may be assigned with different

number of rows to compute but same or similar amount of non-zeros. For simplicity, do not divide a row between two processes but assign a row to a single process only. You do not need to implement OpenMP version of this.

Experimental Study

- You will notice that the serial program we provided requires you to input the number of time steps as the second command line parameter.
- Conduct a strong scaling study on a single node for input file Cube_Coup_dt6.mtx and Flan_1565.mtx with 20 time steps. Observe the running time as you successively double the number of processes, starting at P=1 core and ending at 16 cores, while keeping number of time steps fixed. Compare single processor performance of your parallel MPI code against the performance of the original serial code. Perform this study on Part 1 and Part 3 (if you choose to implement it). You can download the input matrices from the following links:

https://sparse.tamu.edu/MM/Janna/Flan_1565.tar.gz

https://sparse.tamu.edu/MM/Janna/Cube_Coup_dt6.tar.gz

- Repeat single node performance study on MPI+OpenMP parallel code (the code of the second scheme) by using 1, 2, 4, 8, and 16 OpenMP threads per process (thus 16, 8, 4, 2, and 1 MPI processes). In a figure, compare the performance numbers with the strong scaling with MPI-only. Which combination of MPI+OpenMP is the fastest?

Note: For mixing OpenMP and MPI, you will request the same number of total processors and then vary OpenMP threads and MPI processes to span over those processors. For instance, when you request 16 processors, you can use them to spawn 16 processes, 1 OpenMP thread, 8 MPI processes and 2 OpenMP threads and so on.

Grading

Part I: 70 pts (50 pts implementation + 20 pts performance study)

Part II: 20 pts (10 pts implementation + 10 pts performance study)

Part III: 20 pts (optional - bonus)

Report: 10 pts.

You may lose points both for correctness (e.g. deadlock) and performance bugs (e.g. unnecessary synch. point) in your implementation.

Submission

- Document your work in a well-written report which discusses your findings.
- Your report should present a clear evaluation of the design of your code, including bottlenecks of the implementation, and describe any special coding or tuned parameters.

- We have provided a timer to allow you to measure the running time of your serial code. For the MPI code, you could use `MPI_Wtime()` function to collect the execution time of your program (source: <https://www.mcs.anl.gov/research/projects/mpi/tutorial/gropp/node139.html>). Plot both execution time and speedup of your parallel code based on the execution time recorded by this function.
- Submit both the report and source code electronically through blackboard.
- Please create a parent folder named after your username(s). Your parent folder should include a report in pdf and three subdirectories: one for each task. Include all the necessary files to compile your code. Be sure to delete all objects and executable files before creating a zip file.
- GOOD LUCK.

Environment

We will be using KUACC cluster for performance studies. However, if you want to develop and test the application on your local machine:

- You will need an MPI library to be installed on your computer. One MPI implementation is available here: <http://www.open-mpi.org/>. There are other options available. It doesn't matter which implementation you use.
- In order to compile with MPI on your local machine, you need to export paths on your shell or on the development environment (e.g. eclipse). On Unix-based systems, you should add the following lines to your `.bashrc` file:

```
1      # If you wish to use intel compiler
2      source <Path to Intel parallel studio bin folder>/compilervars.sh intel64
3      # MPI paths
4      export PATH=<Path to MPI bin folder>:$PATH
5      export LD_LIBRARY_PATH=<Path to MPI lib folder>:$LD_LIBRARY_PATH
6      export PATH=<Path to MPI include folder>:$PATH
```

- To run a program with MPI, the command line should be as follow.

```
1 mpirun -np <number-of-processes> <path-to-executable>/<executable-name> <arguments>
```

The following list shows the information that you need to use KUACC cluster.

- Accessing KUACC outside of campus requires VPN. You can install VPN through this link: <https://my.ku.edu.tr/faydali-linkler/>
- A detailed explanation is provided in <http://login.kuacc.ku.edu.tr/> to run programs in the KUACC cluster. In this document, we briefly explain it for the Unix-based systems. For other platforms you can refer to the above link.

- In order to log in to the KUACC cluster, you can use ssh (Secure Shell) in a command line as follows: The user name and passwords are the same as your email account.

```
1 bash$ ssh <$username>@login.kuacc.ku.edu.tr
2 bash$ ssh dunat@login.kuacc.ku.edu.tr //example
```

- The machine you logged into is called login node or front-end node. **You are not supposed to run jobs in the login node** but only compile them at the login node. The jobs run on the compute nodes by submitting job scripts.
- To run jobs in the cluster, you have to change your directory to your scratch folder and work from there. The path to your scratch folder is

```
1 bash$ cd /scratch/users/username/
```

- To submit a job to the cluster, you can create and run a shell script with the following command:

```
1 bash$ sbatch <scriptname>.sh
```

- To check the status of your currently running job, you can run the following command:

```
1 bash$ squeue -u <your-user-name>
```

A sample of the shell script is provided in Blackboard along with the assignment folder. In the website of the KUACC cluster, a lot more details are provided.

- To copy any file from your local machine to the cluster, you can use the scp (secure copy) command on your local machine as follows:

```
1 scp -r <filename> <username>@login.kuacc.ku.edu.tr:/kuacc/users/<username>/
2 scp -r src_folder dunat@login.kuacc.ku.edu.tr:/kuacc/users/dunat/ //example
```

-r stands for recursive, so it will copy the src_folder with its subfolders.

- Likewise, in order to copy files from the cluster into the current directory in your local machine, you can use the following command on your local machine:

```
1 scp -r <username>@login.kuacc.ku.edu.tr:/kuacc/users/<username>/fileToBeCopied ./
2 scp -r dunat@login.kuacc.ku.edu.tr:/kuacc/users/dunat/src_code ./ //example
```

- To compile the assignment on the cluster, you can use the MPICH compiler. The compilation commands and flags for the compilers are provided in a Makefile in the assignment folder. Before using the compilers, you firstly need to load their module if they are not already loaded as follows:

```
1 bash$ module avail //shows all available modules in KUACC
2 bash$ module list //list currently loaded modules.
3 bash$ module load mpich/3.2.1 //loads Intel compiler
```

- Here is an example job script to be used on KUACC:

```

1  #!/bin/bash
2  #
3  # You should only work under the /scratch/users/<username> directory.
4  #
5  # Example job submission script
6  #
7  # == Resources ==
8  #
9  #SBATCH --job-name=spmv-jobs
10 #SBATCH --nodes=1
11 #SBATCH --ntasks-per-node=16
12 #SBATCH --partition=short
13 #SBATCH --time=00:30:00
14 #SBATCH --output=spmv-job.out
15
16 #####
17 ##### !!! DO NOT EDIT ABOVE THIS LINE !!! #####
18 #####
19 # Set stack size to unlimited
20 echo "Setting stack size to unlimited..."
21 ulimit -s unlimited
22 ulimit -l unlimited
23 ulimit -a
24 echo
25
26 echo "Running Job...!"
27 echo "=====
28 echo "Running compiled binary..."
29
30 #serial version
31 lscpu
32 echo "Serial version..."
33 build/spmv_serial Cube_Coup_dt6/Cube_Coup_dt6.mtx 20
34 build/spmv_serial Flan_1565/Flan_1565.mtx 20
35
36 #parallel version
37 echo "Parallel version with 1 process"
38 mpirun -np 1 build/spmv Cube_Coup_dt6/Cube_Coup_dt6.mtx 20
39
40 echo "Parallel version with 2 processes"
41 mpirun -np 2 build/spmv Cube_Coup_dt6/Cube_Coup_dt6.mtx 20
42
43 echo "Parallel version with 4 processes"
44 mpirun -np 4 build/spmv Cube_Coup_dt6/Cube_Coup_dt6.mtx 20
45
46 echo "Parallel version with 8 processes"
47 mpirun -np 8 build/spmv Cube_Coup_dt6/Cube_Coup_dt6.mtx 20
48
49 echo "Parallel version with 16 processes"
50 mpirun -np 16 build/spmv Cube_Coup_dt6/Cube_Coup_dt6.mtx 20

```

- If you have problems compiling or running jobs on KUACC, first check the website

provided by the KU IT. If you cannot find the solution there, you can always post a question on Blackboard.

- Don't leave the experiments on KUACC to the last minutes of the deadline as the machine gets busy time to time. Note that there are many other people on campus using the cluster.

References

https://en.wikipedia.org/wiki/Sparse_matrix-vector_multiplication